




# Introduction to Machine Learning in R

Nikolay Oskolkov, MRG Group Leader, LIOS, Riga, Latvia  
R course, 05.02.2026



@oskolkov.bsky.social



Personal homepage:  
<https://nikolay-oskolkov.com>

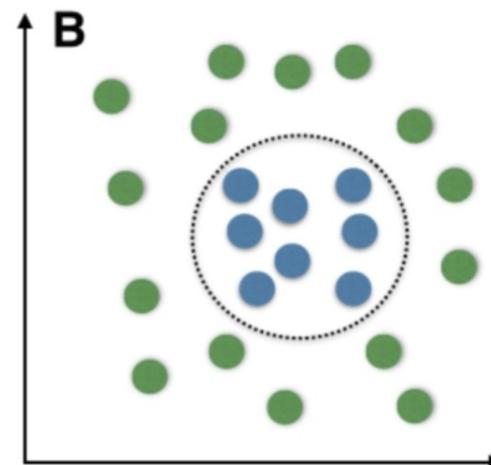
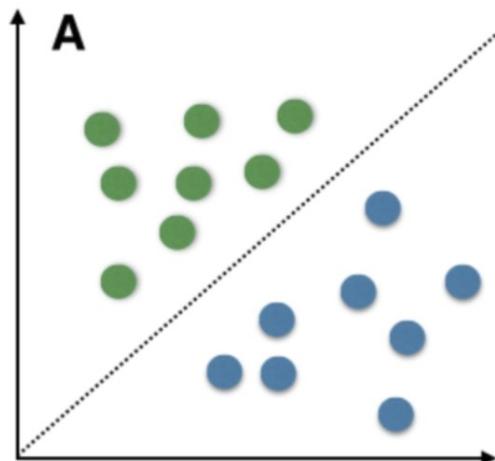
Topics we'll cover in this session:

- 1) How statistics is different from machine learning
- 2) K-means clustering as fingerprint of machine learning
- 3) Basics of machine learning mentality
- 4) Linear machine learning: train, validation and test split
- 5) Overfitting vs. underfitting, coding gradient descent from scratch in R

$Y = f(X)$ , where X is input (data) and Y is output (response)

Y is present – supervised machine learning

Y is absent – unsupervised machine learning



*A: Linearly Separable Data B: Non-Linearly Separable Data*

# Moving from statistics to machine learning

- Statistics is more analytical (pen & paper)

$$L(x_i | \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp^{-\frac{\sum_{i=1}^N (x_i - \mu)^2}{2\sigma^2}}$$

$$\frac{\partial L(x_i | \mu, \sigma^2)}{\partial \mu} = 0; \quad \frac{\partial L(x_i | \mu, \sigma^2)}{\partial \sigma^2} = 0$$

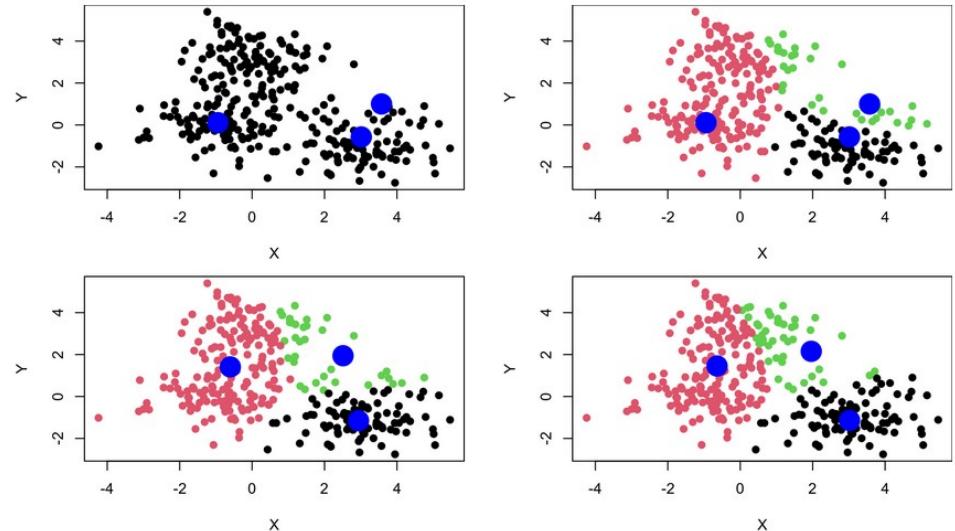
$$\mu = \frac{1}{N} \sum_{i=0}^N x_i - \text{mean estimator}$$

$$\sigma^2 = \frac{1}{N} \sum_{i=0}^N (x_i - \mu)^2 - \text{variance estimator}$$

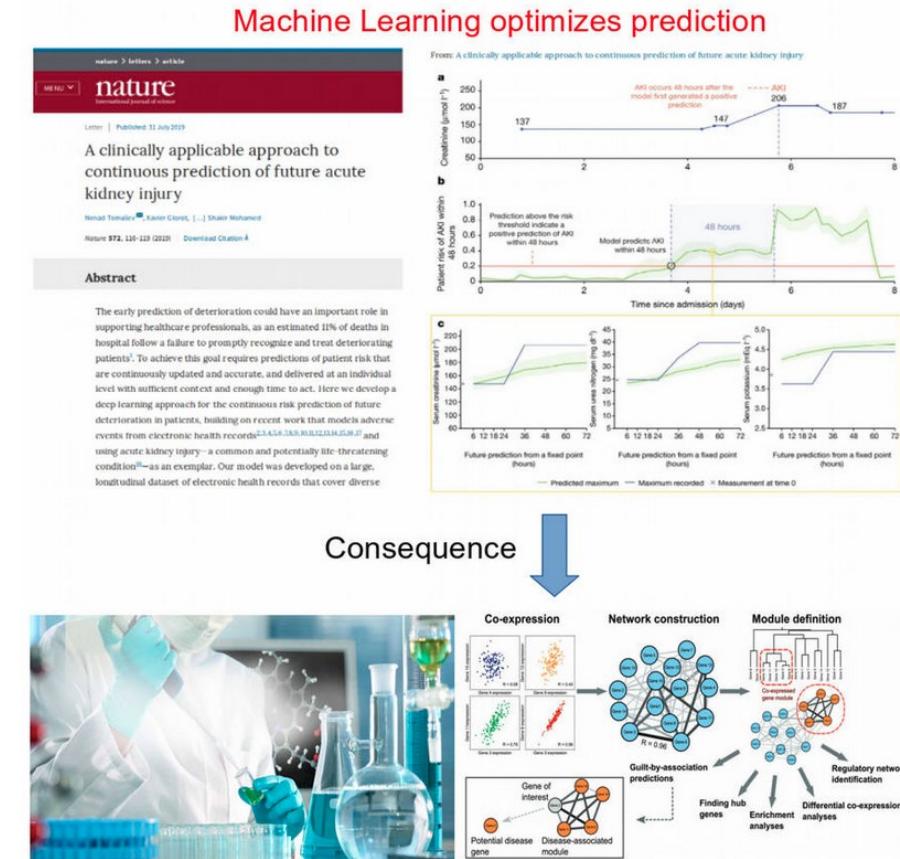
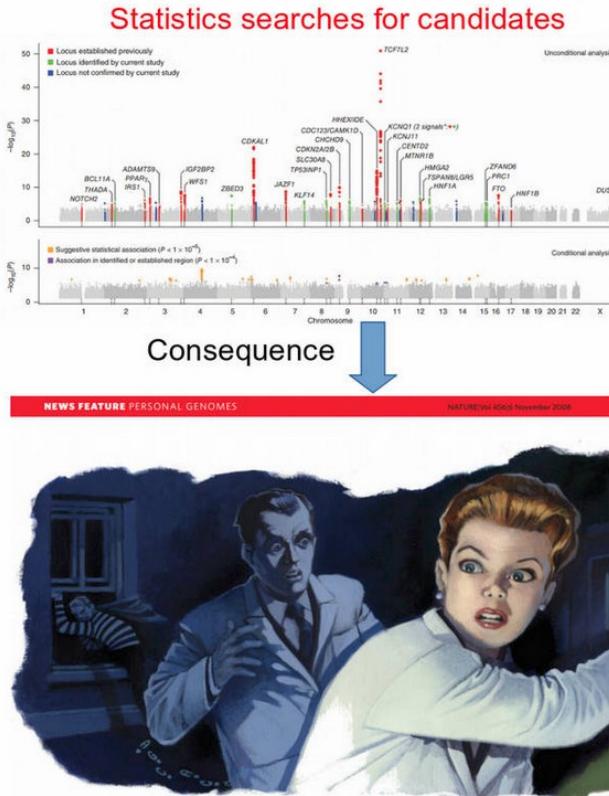
- Machine Learning is more algorithmic (ex. K-means)

```

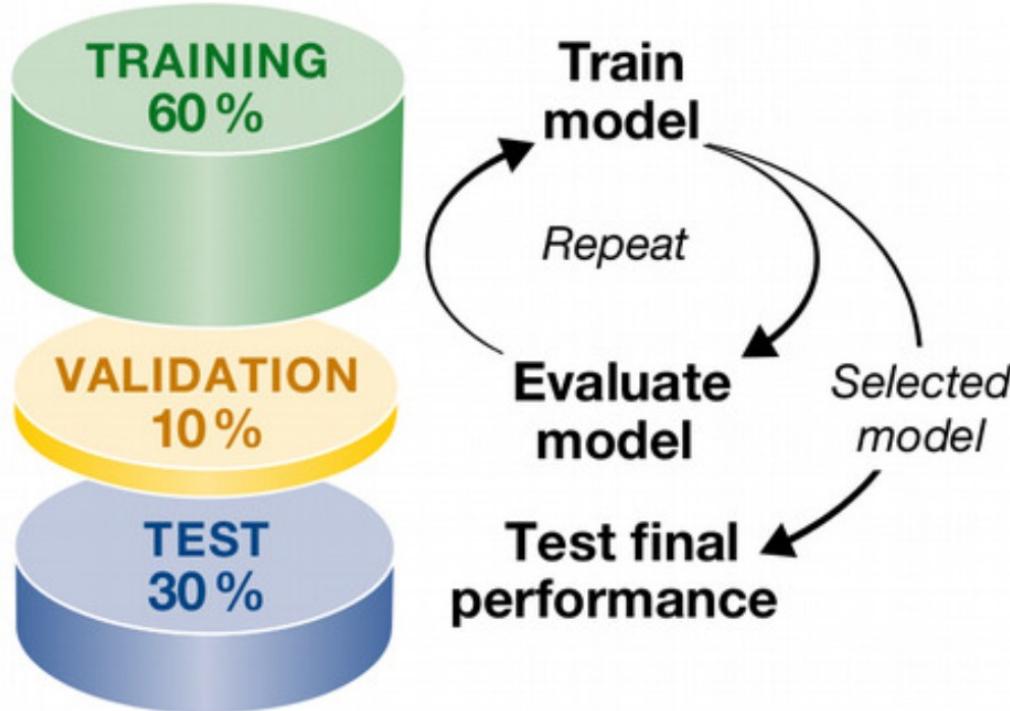
1 K = 3; set.seed(123); c = X[sample(1:dim(X)[1],K),]; par(mfrow=c(2,2),mai=c(0.8,1,0,0))
2 plot(X, xlab = "X", ylab = "Y", pch = 19); points(c, col = "blue", cex = 3, pch = 19)
3 for(t in 1:3)
4 {
5   l <- vector()
6   for(i in 1:dim(X)[1])
7   {
8     d <- vector(); for(j in 1:K){d[j] <- sqrt((X[i,1]-c[j,1])^2 + (X[i,2]-c[j,2])^2)}
9     l[i] <- which.min(d)
10   }
11   plot(X, xlab="X", ylab="Y", col=l, pch=19); points(c, col="blue", cex=3, pch=19)
12   s = list(); for(i in unique(l)){s[[i]] <- colMeans(X[l==i,]); c = Reduce("rbind", s)
13 }
```



# Statistics vs. machine learning: prediction



# How does machine learning work?



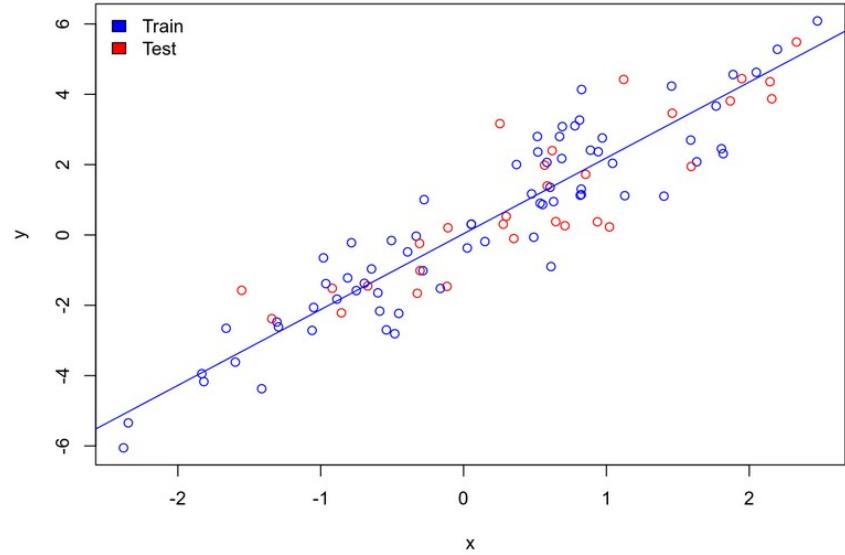
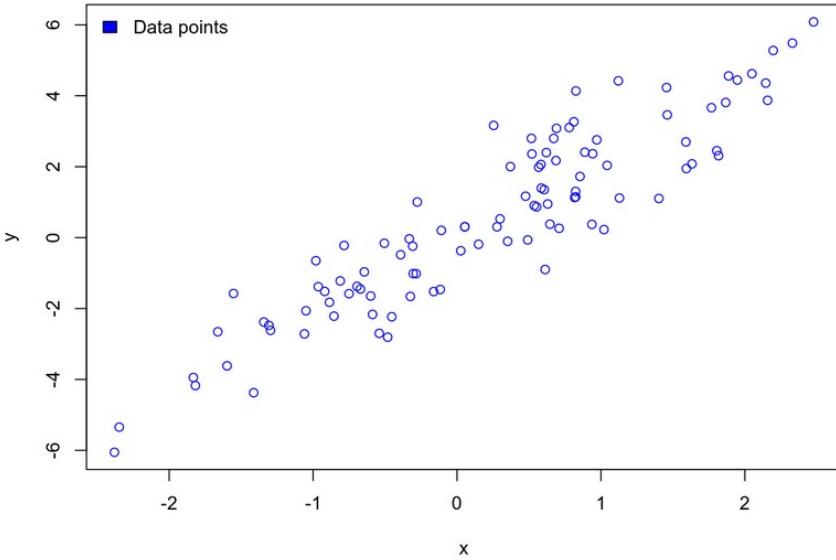
Machine Learning typically involves five basic steps:

1. Split data set into train, validation and test subsets
2. Fit the model on the train subset
3. Validate your model on the validation subset
4. Repeat train - validation split many times and tune hyperparameters
5. Test the accuracy of the optimized model on the test subset.

# Toy example of machine learning

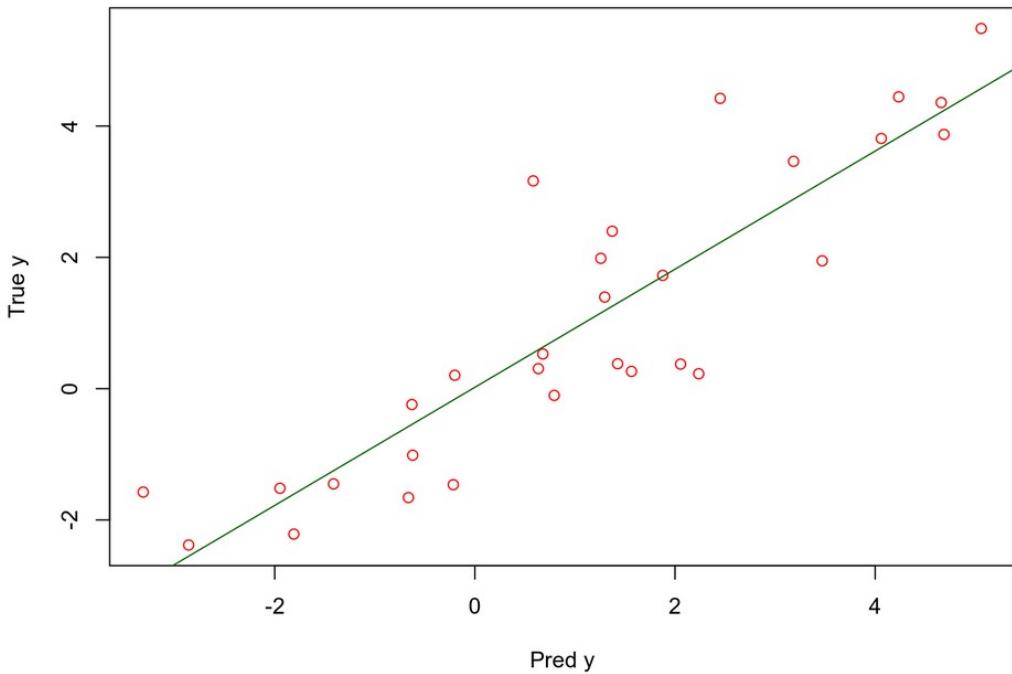
```
1 N <- 100
2 x <- rnorm(N)
3 y <- 2 * x + rnorm(N)
4 df <- data.frame(x, y)
5 plot(y ~ x, data = df, col = "blue")
6 legend("topleft", "Data points", fill = "blue", bty = "n")
```

```
1 train <- df[sample(1:dim(df)[1], 0.7 * dim(df)[1]), ]
2 test <- df[!rownames(df) %in% rownames(train), ]
3 df$col <- ifelse(rownames(df) %in% rownames(test), "red", "blue")
4 plot(y ~ x, data = df, col = df$col)
5 legend("topleft", c("Train", "Test"), fill=c("blue", "red"), bty="n")
6 abline(lm(y ~ x, data = train), col = "blue")
```



# Toy example: model validation

```
1 test_predicted <- as.numeric(predict(lm(y ~ x, data = train), newdata = test))
2 plot(test$y ~ test_predicted, ylab = "True y", xlab = "Pred y", col = "red")
3 abline(lm(test$y ~ test_predicted), col = "darkgreen")
```



```
1 summary(lm(test$y ~ test_predicted))
```

```
Call:
lm(formula = test$y ~ test_predicted)

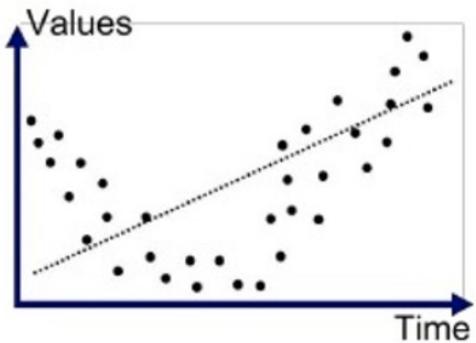
Residuals:
    Min      1Q  Median      3Q     Max 
-1.80597 -0.78005  0.07636  0.52330  2.61924 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept)  0.02058   0.21588   0.095   0.925    
test_predicted 0.89953   0.08678 10.366 4.33e-11 ***  
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1
' ' 1

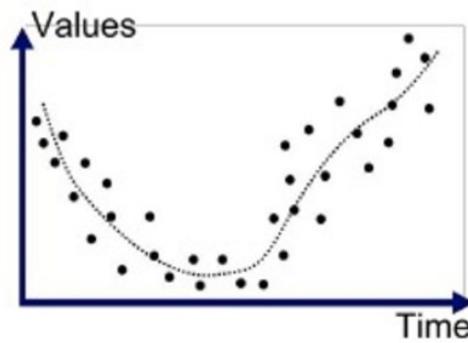
Residual standard error: 1.053 on 28 degrees of freedom
Multiple R-squared:  0.7933,    Adjusted R-squared:
0.7859
F-statistic: 107.4 on 1 and 28 DF,  p-value: 4.329e-11
```

Thus the model explains 79% of variation on the test subset.


## ▼ Underfitting vs. Overfitting



Underfitted



Good Fit/Robust



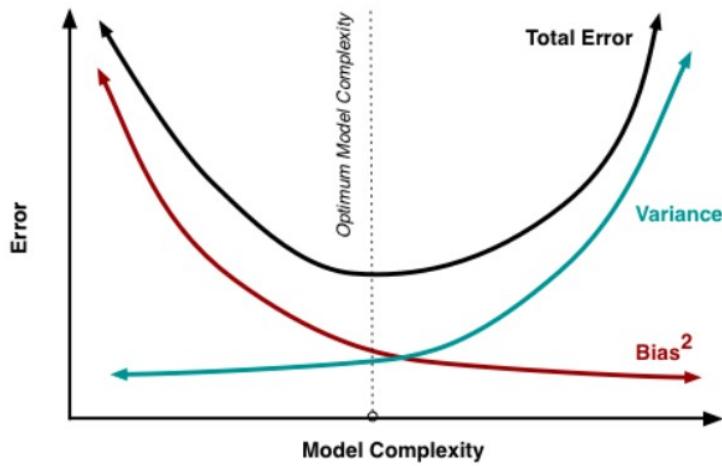
Overfitted

- Akaike Information Criterion (AIC):

$$AIC = 2k - 2\ln(L)$$

- Random Forest: each tree overfitted, but ensemble of trees performs very well

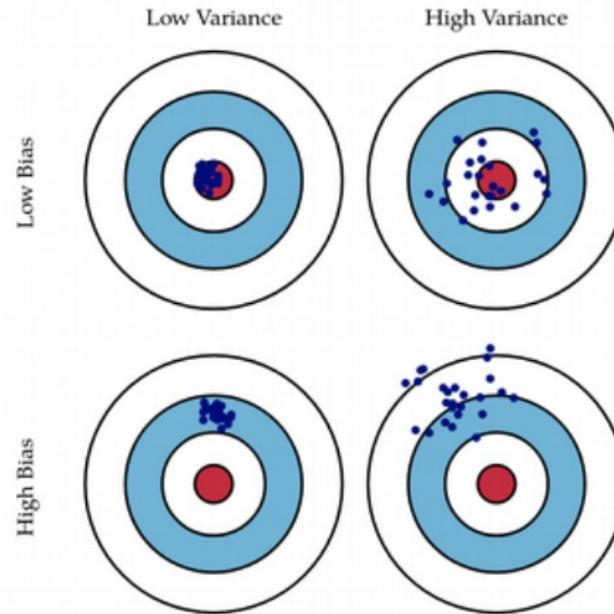
## ▼ Bias-Variance Tradeoff



$$Y = f(X) \implies \text{Reality}$$

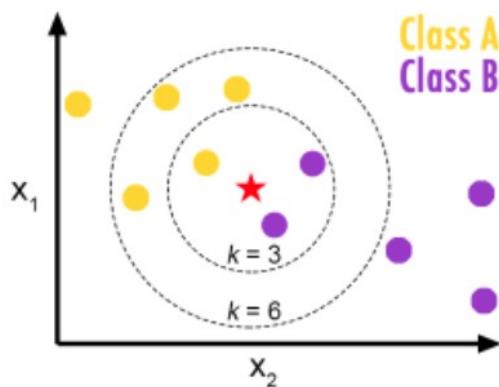
$$\hat{Y} = \hat{f}(X) + \text{Error} \implies \text{Model}$$

$$\text{Error}^2 = (Y - \hat{f}(X))^2 = \text{Bias}^2 + \text{Variance}$$




- It is mathematically proven that ensemble learning keeps the Bias the same but leads to a large decrease of Variance

## ▼ KNN and SVM




- How many out of K neighbors belong to each class
- Majority voting
- Non-linear



- Draw hyperplane that separates classes
- Maximize margins
- Can be linear and non-linear

# Gradient descent



$$y_i = \alpha + \beta x_i + \epsilon, \quad i = 1 \dots n$$

$$E(\alpha, \beta) = \frac{1}{n} \sum_{i=1}^n (y_i - \alpha - \beta x_i)^2$$

$$\hat{\alpha}, \hat{\beta} = \operatorname{argmin} E(\alpha, \beta)$$

$$\frac{\partial E(\alpha, \beta)}{\partial \alpha} = -\frac{2}{n} \sum_{i=1}^n (y_i - \alpha - \beta x_i)$$

$$\frac{\partial E(\alpha, \beta)}{\partial \beta} = -\frac{2}{n} \sum_{i=1}^n x_i (y_i - \alpha - \beta x_i)$$

Numeric implementation of gradient descent:

$$\alpha_{i+1} = \alpha_i - \eta \frac{\partial E(\alpha, \beta)}{\partial \alpha} \Big|_{\alpha=\alpha_i, \beta=\beta_i}$$

$$\beta_{i+1} = \beta_i - \eta \frac{\partial E(\alpha, \beta)}{\partial \beta} \Big|_{\alpha=\alpha_i, \beta=\beta_i}$$

# Coding gradient descent from scratch in R

```
1 n <- 100 # sample size
2 x <- rnorm(n) # simulated explanatory variable
3 y <- 3 + 2 * x + rnorm(n) # simulated response variable
4 summary(lm(y ~ x))
```

Call:  
lm(formula = y ~ x)

Residuals:

Min	1Q	Median	3Q	Max
-1.9073	-0.6835	-0.0875	0.5806	3.2904

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	2.89720	0.09755	29.70	<2e-16 ***
x	1.94753	0.10688	18.22	<2e-16 ***

---


Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.9707 on 98 degrees of freedom  
Multiple R-squared: 0.7721, Adjusted R-squared: 0.7698  
F-statistic: 332 on 1 and 98 DF, p-value: < 2.2e-16

Let us now reconstruct the intercept and slope from gradient descent

```
1 alpha <- vector(); beta <- vector()
2 E <- vector(); dEdalpha <- vector(); dEdbeta <- vector()
3 eta <- 0.01; alpha[1] <- 1; beta[1] <- 1 # initialize alpha and beta
4 for(i in 1:1000)
5 {
6   E[i] <- (1/n) * sum((y - alpha[i] - beta[i] * x)^2)
7   dEdalpha[i] <- -sum(2 * (y - alpha[i] - beta[i] * x)) / n
8   dEdbeta[i] <- -sum(2 * x * (y - alpha[i] - beta[i] * x)) / n
9
10  alpha[i+1] <- alpha[i] - eta * dEdalpha[i]
11  beta[i+1] <- beta[i] - eta * dEdbeta[i]
12 }
13 print(paste0("alpha = ", tail(alpha, 1), ", beta = ", tail(beta, 1)))
```

```
[1] "alpha = 2.89719694937354, beta = 1.94752837381973"
```



Take home messages of the session:

- 1) Machine Learning is different from statistics in terms of more algorithmic approach and prediction
- 2) Train, validation and test are fundamental machine learning steps that aim at improving model generalizability
- 3) Challenge in machine learning is to find a balance between overfitting and underfitting
- 4) Gradient descent is the core engine of machine learning and is based on computing derivatives of the loss function



# Acknowledgments: LIOS + TARGETWISE



2027  
National  
Development Plan



Funded by  
the European Union